

Moby Dicks

*The reading app that helps people
actually finish the books they start.*

AATISH · BRANDT · SAI

Product pitch · 4 minutes

Reading isn't a discovery problem.

12.6

books a year started

—

5

books a year finished

More than half of US adults didn't finish a single book last year.

Reddit, X, and book forums are full of the same question — "why can't I finish books anymore?" People aren't short on recommendations. They're short on follow-through.



Everyone optimizes for the wrong thing.

Goodreads

150M users

OPTIMIZES FOR

What you'll like

THE GAP

Star ratings + "want to read" lists. No follow-through, no completion data, no help finishing.

ChatGPT

general LLM

OPTIMIZES FOR

An 8-message chat

THE GAP

Generic recommendations, hallucinated titles, no memory of what you've actually read or abandoned.

Blinkist · Headway

\$50M / \$160M ARR

OPTIMIZES FOR

Skip the book entirely

THE GAP

15-minute summaries. Solves the problem by not reading at all — different product, different user.

Nobody is building for completion. We are.

A coach, not a list.

MOBY DICKS

TOPIC

Japanese magical realism

TIME

Evening read

DIFFICULTY

Moderate

Generate my curriculum →

Returns in ~10 seconds

HOW IT WORKS

01 *One survey*

Topic, time, difficulty. The whole onboarding fits on one screen.

02 *Curated curriculum*

Five real books, ordered by progressive difficulty, with a justification per pick.

03 *Check-ins after every book*

"Too easy, too dense, just right?" We re-rank the rest of your plan from your answer.



A real curriculum from a real query.

INPUT

Japanese magical realism · evening read · moderate

WEEK 1

***Kafka on the
Shore****Haruki Murakami*

WEEK 2

***The Wind-Up
Bird Chronicle****Haruki Murakami*

WEEK 3

***Strange
Weather in
Tokyo****Hiromi Kawakami*

WEEK 4

***The Memory
Police****Yoko Ogawa*

WEEK 5

1Q84*Haruki Murakami*

10 seconds end-to-end · **\$0.001** per curriculum · **0 hallucinations** — every pick is a real book

Big market, perfect moment.

TAM

190M

US adults who read at least one book per year

75% of US adults · Pew Research

SAM

75M

US readers who started but didn't finish a book last year

WordsRated · Book Riot, 2023

SOM (Y1-3)

\$31M

ARR target — 0.5% of SAM × \$84/yr

375K paying users at targeted conversion

WHY NOW

1. Per-user LLM curation just got cheap. Gemini Flash dropped per-request cost ~25× since 2023.
2. BookTok and Bookstagram have made reading a mainstream identity again — but completion rates keep falling.
3. Reading-app subscriptions work. Headway hit \$160M revenue in 2025; Blinkist crossed 26M users.

Freemium with 95%+ gross margin.

FREE

\$0

Acquisition tier

- 1 curriculum per month
- Read-only progress view
- Affiliate links to bookstore.org / Bookshop

Top of funnel · drives word-of-mouth

PRO

\$7 / mo

\$60 / year — save 30%

- Unlimited curricula
- Check-ins + curriculum reoptimization
- Momentum dashboard + reading history
- Social: shared curricula with friends

Target: 0.5%–2% of free users convert

UNIT ECONOMICS

\$84

Pro ARPU / year

\$0.36

LLM cost per user / year

~95%

Gross margin at scale

Day 1

Profitable from user one



Distribution where readers already live.

GO-TO-MARKET

PHASE 1 · NOW

Reading communities

r/books, r/suggestmeabook, BookTok creators, Substack reading newsletters.
Low CAC, high intent.

PHASE 2 · Y1-2

SEO content engine

"What to read after [book]" landing pages. Goodreads SEO is broken — we fill the gap.

PHASE 3 · Y2+

B2B2C licensing

White-label for libraries, indie bookstores, and book clubs. ABA partnership in target.

DEFENSIBLE MOAT

Behavioral completion data

Every check-in is a labeled training example. Nobody else has this.

Compounding personalization

Each curriculum gets smarter from your prior reads — switching cost grows.

Cost structure

\$0.001 per curriculum means we're profitable from user one.

First mover on completion

Goodreads can't pivot here without breaking their core ratings model.

VISION

Strava for readers.

Every reader gets a coach. Coaches turn starters into finishers.

TRY IT

github.com/SaiMulagan/mobydick

AATISH · BRANDT · SAI Thank you.

Post-pivot reality: \$8/month at MVP scale.

PROJECTED MONTHLY

\$7.91

at 50 users · 3 calls each

80% cheaper than the Milestone 1 plan (\$40.07).

\$0.00053 per request

MONTHLY LINE ITEMS · 150 requests / month

Component	Then (M1)	Now (M6)
Cloud Run (scale-to-zero)	\$5.00	\$5.00
Cloud SQL → BigQuery	\$30.00	\$2.40
Operations / Logging	\$5.00	\$0.50
Vertex AI · Gemini Flash	\$0.07	\$0.01
Total	\$40.07	\$7.91

Annual run-rate: ~\$95. Comfortably inside the \$150 GCP credit cap.

Why these numbers, not the old ones.

01 *Cloud SQL → BigQuery*

\$30.00 → \$2.40 / month

Original plan needed Postgres for relational user/curriculum data. Once we realized the candidate scan is a one-shot read-only query against a static catalog, BigQuery on-demand became the right answer. We pay only for bytes scanned (≈400 MB / 150 queries) and storage (≈12 GB at \$0.02/GB).

02 *GPT-4o → Gemini Flash*

~25× cheaper per token

\$0.075 / M input tokens vs \$5 / M for GPT-4o. Co-located with Cloud Run on GCP, so auth is IAM (no rotating keys) and latency stays under a second. Quality difference is invisible for our schema-constrained JSON output.

03 *Scale-to-zero Cloud Run*

\$0 idle cost

Min instances = 0 in dev, 1 only on demo day. A cold start adds ~2 s the first time, hidden behind our existing 10 s curriculum-generation latency. No always-on container means our infra cost tracks usage, not calendar months.

04 *Two-call pipeline + caching*

\$0.00053 / request

Splitting extraction from curation lets us cache by (topic + profile_bucket) in the curricula table, so the second user with the same query pays \$0. With 50 users and natural query overlap, our actual LLM spend is well under our worst-case estimate.

Net result: Profitable from user one. Nothing on this list requires further engineering.